

2 Work-Stealing for Task Trees — extended¹

Work stealing is a popular efficient technique for performing load balancing in multicore computations. In traditional schemes, the work-stealing is *receiver-initiated*: workers that run out of work are responsible for *stealing* tasks. In a dual approach, called *sender-initiated work-stealing*, workers with tasks are responsible for actively sharing their tasks with workers that are out of work.

In this challenge we investigate work-stealing in the context of a *binary tree of tasks*. Each task has at most two child *subtasks* and a task may be executed only *after* parent task; the root task must be executed first, therefore. The order in which tasks can be executed is otherwise not restricted.

Let $P > 0$ be the number of workers; each worker has a unique ID i in the range 0 to $P-1$. Each worker has a double-ended queue $q[i]$ representing the tasks currently ready to execute and assigned to this worker; these start off empty, and the root task is then assigned to worker 0. Our algorithms are each built around the same key main function, shown in pseudo code below:

```
typedef int task          // tasks are represented by their IDs
const int nTasks          // number of tasks (IDs 0,1,...,nTasks-1)
const task NO_TASK = -1 // special code to denote 'no child task'
const task ROOT_TASK = 0 // ID of the root task
task subtask[nTasks][2] // maps task ID to child task IDs / NO_TASK
bool executed[nTasks]   // marks executed tasks; initially 0 (false)

const int P              // number of workers
deque<task> q[P]          // double-ended queue per worker; initially empty

// entry point for each worker (calling this concurrently)
void main(int i) // i = ID of the worker; 0 for the initial worker
  if i = 0 // the worker who starts things off
    push_bottom(q[i], ROOT_TASK)
  repeat // until termination
    if (empty(q[i])) // if out of work, try to acquire a task
      acquire(i)     // scheme-dependent function; see later
    else // pick a task and execute it
      task t = pop_bottom(q[i])
      communicate(i) // scheme-dependent function; see later
      execute(i, t)

// execution of a task t by worker i
void execute(int i, task t)
  // perform some task-specific computation; omitted for simplicity
  executed[t] += 1 // flag the task as executed
  // then schedule the subtasks
  add_task(i, subtask[t][1])
  add_task(i, subtask[t][0])

// called for scheduling a task t into worker i's queue
void add_task(int i, task t)
  if t != NO_TASK
    push_bottom(q[i], t)
```

¹We warmly thank Arthur Charguéraud for contributing the idea for this challenge.

The array `subtask` (which is never mutated in the code; you may assume this to be immutable if it helps you) expresses the tree structure of tasks: looking up a task's ID in the array gives an array with two elements, storing the task IDs of its respective subtasks (or the special value `NO_TASK`).

We assume an existing suitable implementation of a double-ended queue (the `deque` type in our pseudo code). You do *not* need to implement this type for these challenges. However, since the code we are concerned with interacts with these queues, you will need specifications for five functions on these queues:

`empty(Q)` returning a boolean indicating whether the queue `Q` is empty.

`peek_top(Q)` returning the element at the start of `Q` (without removing it).

`pop_top(Q)` removing the top element from `Q` and returning it.

`push_bottom(Q,T)` which modifies the queue `Q`, adding the task `T` at the end.

`pop_bottom(Q)` removing the bottom/end element from `Q` and returning it.

The variation between different task-handling schemes is expressed by changing (only) the implementations of the `acquire` and `communicate` functions.

Version 0: Sequential task processing We start with a sequential scheme: here, we can assume $P = 1$ and so there is a unique worker executing the `main` function with ID (`i` parameter) 0. In this version, the two functions `acquire` and `communicate` are no-ops: their implementations are empty, and calling them does nothing. The (only) worker will initially add the root task to its queue, and continually execute a task in its queue, queueing up its subtasks, and so on. We do not handle *worker* termination in the code (which is complex for concurrent schemes), but all *tasks* should eventually be executed this way.

Tasks for version 0

- (a) Formalise the assumption that the initial values stored in the array (of length two arrays) `subtask` define a valid binary tree rooted at task ID 0.
- (b) Define suitable specifications for the queue functions listed above.
- (c) Verify that the pseudocode functions given are memory-safe / crash free (assuming that the queue functions are similarly safe): in particular, verify that all array accesses performed are guaranteed to be within bounds.
- (d) Verify that every task is executed at most once.
- (e) Verify that all task dependencies (as expressed by the tree structure) are respected: a subtask is never executed before its parent.
- (f) Verify that all tasks are eventually executed.
- (g) Verify that (after `ROOT_TASK` has been inserted by worker 0) all the worker queues (in version 0, the single queue `q[0]`) eventually become empty.

Version 1: Sender-initiated Work-stealing The following alternative implementations of the `acquire` and `communicate` functions (along with additional definitions/state as shown) implement a *sender-initiated work-stealing* scheme.

```
// extra definitions/state
const task WAITING = -2 // code for 'a task would be welcome'
const task NOT_WAITING = -3 // code for 'not receiving tasks'
task s[P]; // communication cells, all initially NOT_WAITING

// called by workers when running out of work
void acquire(int i)
    s[i] = WAITING // 'a task would be welcome'
    while (s[i] == WAITING) // block until receiving a task
        noop
    add_task(i, s[i])
    s[i] = NOT_WAITING // technically optional

// consider pushing a task to an idle (different) worker
void communicate(int i)
    if (empty(q[i])) // cannot provide a task if we have none
        return
    int j = random in {0, ..., P-1}\{i} // pick a random other worker
    if s[j] != WAITING // check if that worker is waiting for tasks
        return // give up if we picked a worker not waiting
    task t = peek_top(q[i])
    // attempt to atomically take the target communication slot
    bool r = compare_and_swap(&s[j], WAITING, t)
    if (r) // we successfully wrote the task ID to s[j]
        pop_top(q[i]) // remove from our queue; now worker j gets it
```

The `acquire` function is only called when a worker has no tasks, and sets the worker's communication cell to `WAITING` to signal that a task can be assigned to it. It then busy-waits until this cell's value has been changed (to a task ID), and it then inserts this task and continues.

The `communicate` function represents worker `i` considering sending a task to another worker. If it has a task to send, it randomly guesses (once) a different worker ID, and checks whether that worker is waiting for work (if not, it gives up for on communication for now). If so, it uses a compare-and-swap operation (which might be racing with other workers trying to assign the same worker a task) to attempt to atomically update the worker's communication cell with the task ID to send, removing this from its own queue if this operation is successful.

Tasks for version 1

- (h) Prove that all functions provided for this scheme are memory-safe / crash free (as for task (c) above).
- (i) Prove that the same properties (d)–(g) as for version 0 for this new sender-initiated concurrent scheme hold (in particular, assuming $P > 1$).
- (j) Write a short textual comment labelled **MODULARITY**: explaining to what extent you are able to reuse parts of the *code* and *verification effort/results* between your two different versions.

NEW CONTENT FROM HERE!

Version 2: Receiver-initiated Work-stealing The following alternative implementations of the `acquire` and `communicate` functions (along with definitions/state as shown) implement a *receiver-initiated work-stealing* scheme.

```
const int NO_REQUEST = -1 // special "worker ID" value
int r[P] // request cell per worker; initially all NO_REQUEST
const task NO_RESPONSE = -2 // code for 'no task provided yet'
task t[P] // transfer cell per worker; initially all NO_RESPONSE

// called by workers when running out of work
void acquire(int i)
    while true // block until receiving a proper task
        t[i] = NO_RESPONSE // initialize the cell for receiving a task
        int k = random in {0, .., P-1}\{i} // pick random other worker
        if compare_and_swap(&r[k], NO_REQUEST, i) // make a request
            while (t[i] == NO_RESPONSE) // wait for a response
                communicate(i) // reply negatively to incoming queries
            if (t[i] != NO_TASK) // if we obtained a valid task
                add_task(i, t[i]) // get ready to work on that task
            return
        // otherwise, if obtained a negative reply, then try again
        communicate(i) // provide negative reply to incoming queries

// check for incoming steal requests
void communicate(int i)
    int j = r[i] // check our own request cell
    if j == NO_REQUEST // if no request, then nothing to do
        return
    if (empty(q[i]))
        t[j] = NO_TASK // if no task at hand, provide a negative reply
    else
        t[j] = pop_top(q[i]) // else, reply with a task
    r[i] = NO_REQUEST // reset request cell to allow further requests
```

The scheme here is for workers without work to (via `acquire`) first prepare their transfer cell for receiving a task, then pick a random other worker and register a request for work in their request cell. Then workers wait to receive a response: either `NO_TASK` or a task ID (signifying a task transferred to them). All workers are responsible for checking whether they have received requests and responding, by periodically calling the `communicate` function.

Tasks for version 2

- (k) Prove that all functions provided for this scheme are memory-safe / crash free (as for task (c) above).
- (l) Prove that the same properties (d)–(g) as for version 1 (for $P > 1$).
- (m) Write a short textual comment labelled **MODULARITY**: explaining to what extent you are able to reuse parts of the *code* and *verification effort/results* between each of your different versions.